


High Performance Computing

Ch. 2


Basic optimization techniques for serial code

Computer Eng., KMITL
Assoc. Prof. Dr. Surin. K.



Basic optimization techniques for serial code

- 2.1 Scalar profiling
- 2.2 Common sense optimizations
- 2.3 Simple measures, large impact
- 2.4 The role of compilers
- 2.5 C++ optimizations





2.1 Scalar Profiling

1	%	cumulative	self		self	total	
2	time	seconds	seconds	calls	ms/call	ms/call	name
3	70.45	5.14	5.14	26074562	0.00	0.00	intersect
4	26.01	7.03	1.90	4000000	0.00	0.00	shade
5	3.72	7.30	0.27	100	2.71	73.03	calc_tile

% time Percentage of overall program runtime used *exclusively* by this function, i.e., not counting any of its callees.


cumulative seconds Cumulative sum of exclusive runtimes of all functions up to and including this one.





2.1 Scalar Profiling

1	%	cumulative	self		self	total	
2	time	seconds	seconds	calls	ms/call	ms/call	name
3	70.45	5.14	5.14	26074562	0.00	0.00	intersect
4	26.01	7.03	1.90	4000000	0.00	0.00	shade
5	3.72	7.30	0.27	100	2.71	73.03	calc_tile




self seconds Number of seconds used by this function (exclusive). By default, the list is sorted according to this field.

calls The number of times this function was called.

self ms/call Average number of milliseconds per call that were spent in this function (exclusive).


total ms/call Average number of milliseconds per call that were spent in this function, including its callees (inclusive).



2.1 Scalar Profiling

of sampling hits (first column) and the relative percentage of total program samples (second column):

1			:	DO 215 M=1,3
2	4292	0.9317	:	bremsdir(M) = bremsdir(M) + FH(M)*Z12
3	1462	0.3174	:	215 CONTINUE
4			:	
5	682	0.1481	:	U12 = U12 + GCL12 * Upot
6			:	
7			:	DO 230 M=1,3
8	3348	0.7268	:	F(M,I)=F(M,I)+FH(M)*Z12
9	1497	0.3250	:	Fion(M)=Fion(M)+FH(M)*Z12
10	501	0.1088	:	230 CONTINUE



2.1 Scalar Profiling

1	index	% time	self	children	called	name
2			0.27	7.03	100/100	main [2]
3	[1]	99.9	0.27	7.03	100	calc_tile [1]
4			1.90	5.14	4000000/4000000	shade [3]
5	-----					
6						<spontaneous>
7	[2]	99.9	0.00	7.30		main [2]
8			0.27	7.03	100/100	calc_tile [1]
9	-----					
10					5517592	shade [3]
11			1.90	5.14	4000000/4000000	calc_tile [1]
12	[3]	96.2	1.90	5.14	4000000+5517592	shade [3]
13			5.14	0.00	26074562/26074562	intersect [4]
14					5517592	shade [3]
15	-----					
16			5.14	0.00	26074562/26074562	shade [3]
17	[4]	70.2	5.14	0.00	26074562	intersect [4]


2.1 Scalar Profiling

% time The percentage of overall runtime spent in this function, including its callees (inclusive time). This should be identical to the product of the number of calls and the time per call on the flat profile.


self For each indexed function, this is exclusive execution time (identical to flat profile). For its callers (callees), it denotes the inclusive time this function (each callee) contributed to each caller (this function).

children For each indexed function, this is inclusive minus exclusive runtime, i.e., the contribution of all its callees to inclusive time. Part of this time contributes to inclusive runtime of each of the function's callers and is denoted in the respective caller rows. The callee rows in this column designate the contribution of each callee's callees to the function's inclusive runtime.

called denotes the number of times the function was called (probably split into recursive plus nonrecursive contributions, as shown in case of `shade()` above). Which fraction of the number of calls came from each caller is shown in the caller row, whereas the fraction of calls for each callee that was initiated from this function can be found in the callee rows.



Sampling histogram (solid) with
number of samples vs. source code line number.






2.1.2 Hardware performance counters

1	CPU Cycles.....	8721026107
2	Retired Instructions.....	21036052778
3	Average number of retired instructions per cycle.....	2.398151
4	L2 Misses.....	101822
5	Bus Memory Transactions.....	54413
6	Average MB/s requested by L2.....	2.241689
7	Average Bus Bandwidth (MB/s).....	1.197943
8	Retired Loads.....	694058538
9	Retired Stores.....	199529719
10	Retired FP Operations.....	7134186664
11	Average MFLOP/s.....	1225.702566
12	Full Pipe Bubbles in Main Pipe.....	3565110974
13	Percent stall/bubble cycles.....	40.642963



2.1.2 Hardware performance counters

1	CPU Cycles.....	28526301346
2	Retired Instructions.....	15720706664
3	Average number of retired instructions per cycle.....	0.551095
4	L2 Misses.....	605101189
5	Bus Memory Transactions.....	751366092
6	Average MB/s requested by L2.....	4058.535901
7	Average Bus Bandwidth (MB/s).....	5028.015243
8	Retired Loads.....	3756854692
9	Retired Stores.....	2472009027
10	Retired FP Operations.....	4800014764
11	Average MFLOP/s.....	252.399428
12	Full Pipe Bubbles in Main Pipe.....	25550004147
13	Percent stall/bubble cycles.....	89.566481




2.2 Common sense optimizations: Do less work!

```
1 logical :: FLAG
2 FLAG = .false.
3 do i=1,N
4   if(complex_func(A(i)) < THRESHOLD) then
5     FLAG = .true.
6   endif
7 enddo
```

If `complex_func()` has no side effects, the only information that gets communicated to the outside of the loop is the value of `FLAG`. In this case, depending on the probability for the conditional to be true, much computational effort can be saved by leaving the loop as soon as `FLAG` changes state:


```
1 logical :: FLAG
2 FLAG = .false.
3 do i=1,N
4   if(complex_func(A(i)) < THRESHOLD) then
5     FLAG = .true.
6     exit
7   endif
8 enddo
```




2.3.2 Avoiding branches

```
1 do j=1,N
2   do i=1,N
3     if(i.ge.j) then
4       sign=1.d0
5     else if(i.lt.j) then
6       sign=-1.d0
7     else
8       sign=0.d0
9     endif
10    C(j) = C(j) + sign * A(i,j) * B(
11  enddo
12 enddo
```

```
1 do j=1,N
2   do i=j+1,N
3     C(j) = C(j) + A(i,j) * B(i)
4   enddo
5 enddo
6 do j=1,N
7   do i=1,j-1
8     C(j) = C(j) - A(i,j) * B(i)
9   enddo
10 enddo
```





2.3.3 Using SIMD instruction sets

```
1  ! vectorized part
2  rest = mod(N,4)
3  do i=1,N-rest,4
4      load R1 = [x(i),x(i+1),x(i+2),x(i+3)]
5      load R2 = [y(i),y(i+1),y(i+2),y(i+3)]
6      ! "packed" addition (4 SP flops)
7      R3 = ADD(R1,R2)
8      store [r(i),r(i+1),r(i+2),r(i+3)] = R3
9  enddo
10 ! remainder loop
11 do i=N-rest+1,N
12     r(i) = x(i) + y(i)
13 enddo
```

```
1  real, dimension(1:N) :: r, x, y
2  do i=1, N
3      r(i) = x(i) + y(i)
4  enddo
```

