


High Performance Computing

Shared Memory Parallel CPU & OpenMP


Computer Eng., KMITL
Assoc. Prof. Dr. Surin. K.

The good old days ...

Moore's Law



From Hennessy and Patterson, *Computer Architecture: A Quantitative Approach*, 4th edition, Sept. 15, 2006



Moore's Law






Figure 3.2: Progress of maximum arithmetic performance (open circles) and peak theoretical memory bandwidth (filled circles) for Intel processors since 1994. The fastest processor in terms of clock frequency is shown for each year. (Data collected by Jan Treibig.)

Number of systems versus core count




Microprocessor trends


Individual processors are many core (and often heterogeneous) processors.




Intel SCC Processor




AMD ATI RV770



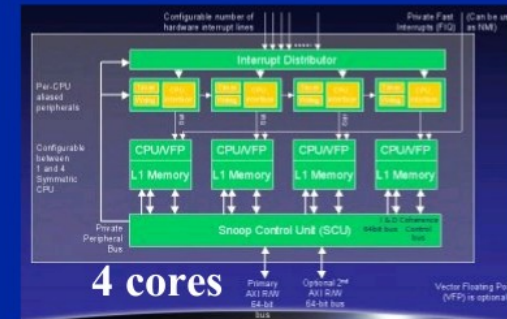
NVIDIA Tesla C1060



Intel® Xeon® processor




IBM Cell




ARM MPCORE

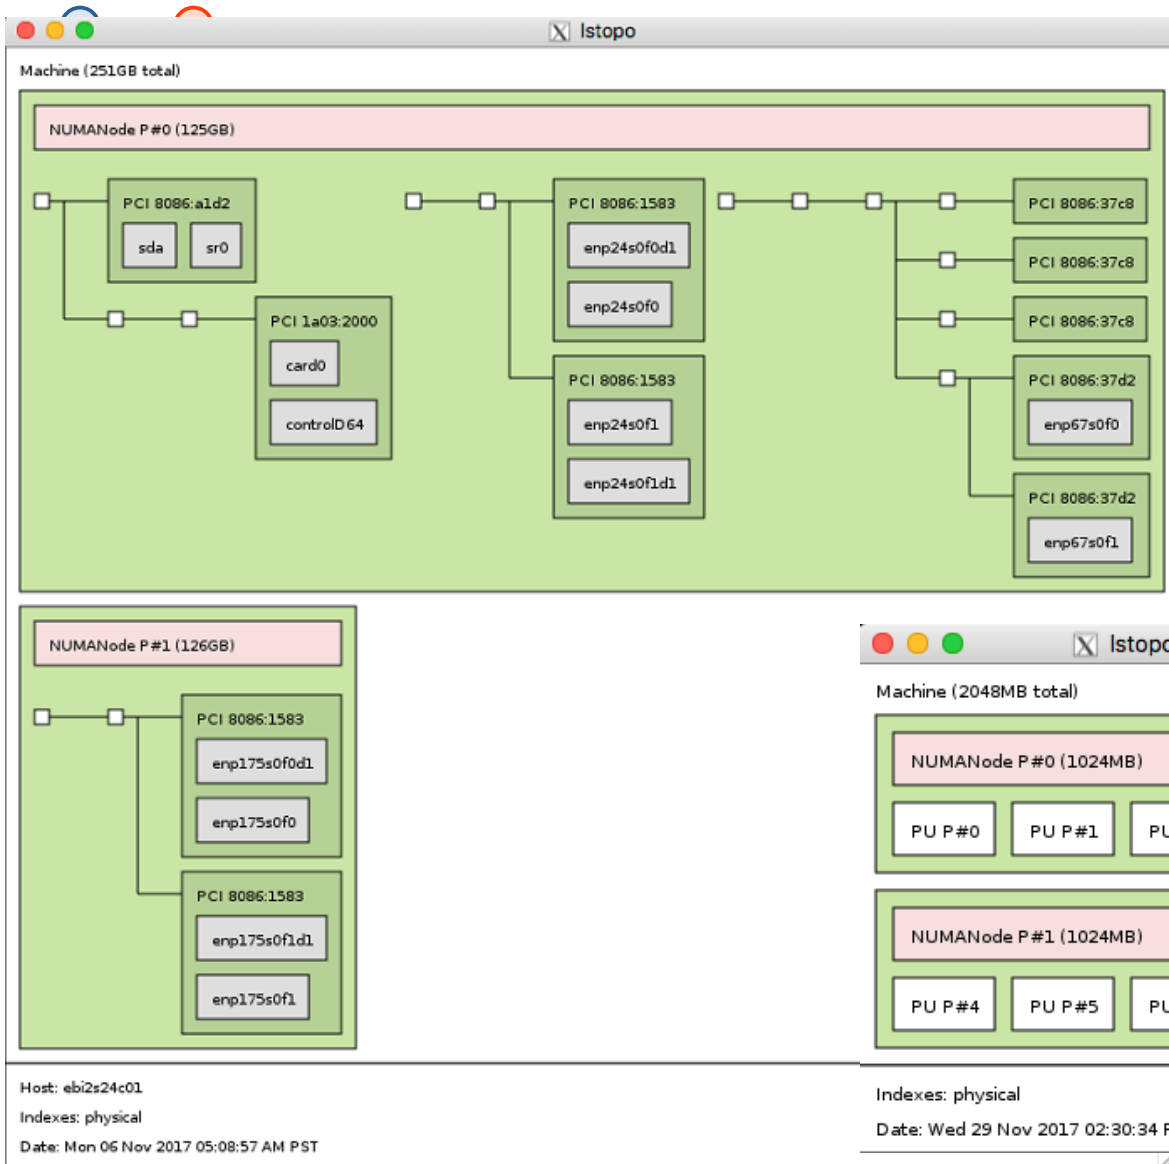
NUMA: Intel Sandy Bridge

CPU architecture (Intel Sandy Bridge)

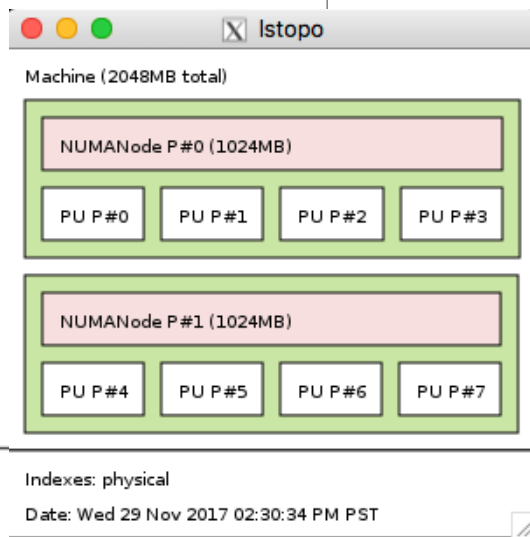


NUMA: Intel Quick Assist Technology



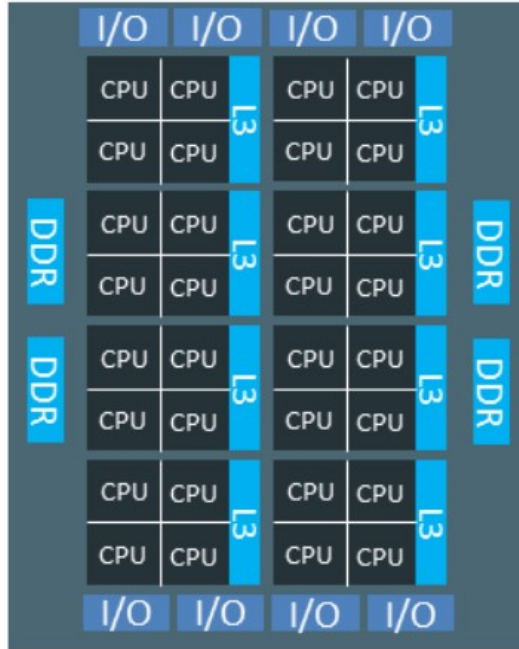


NUMA: Intel Quick Assist Technology



NUMA: AMD Epyc

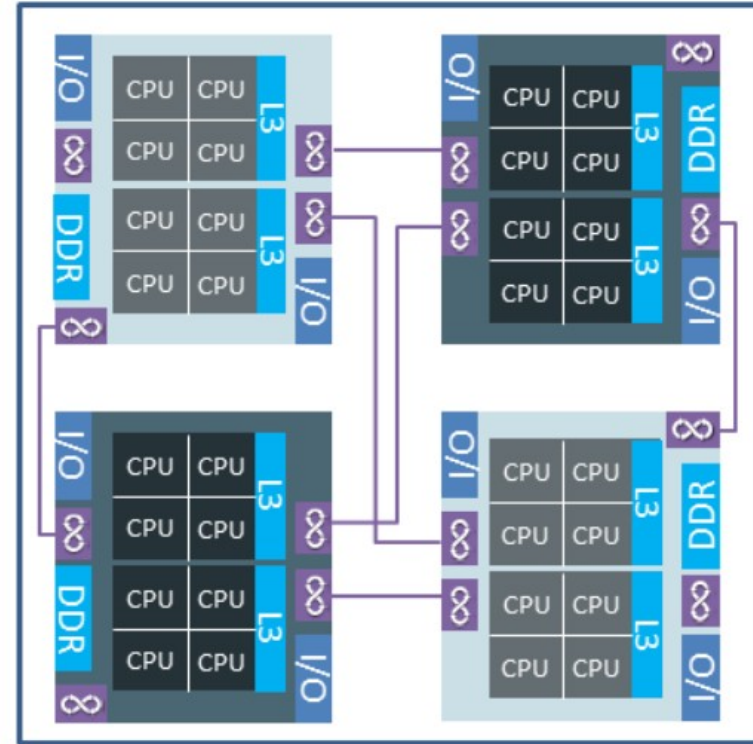
Monolithic Die



32C Die Cost

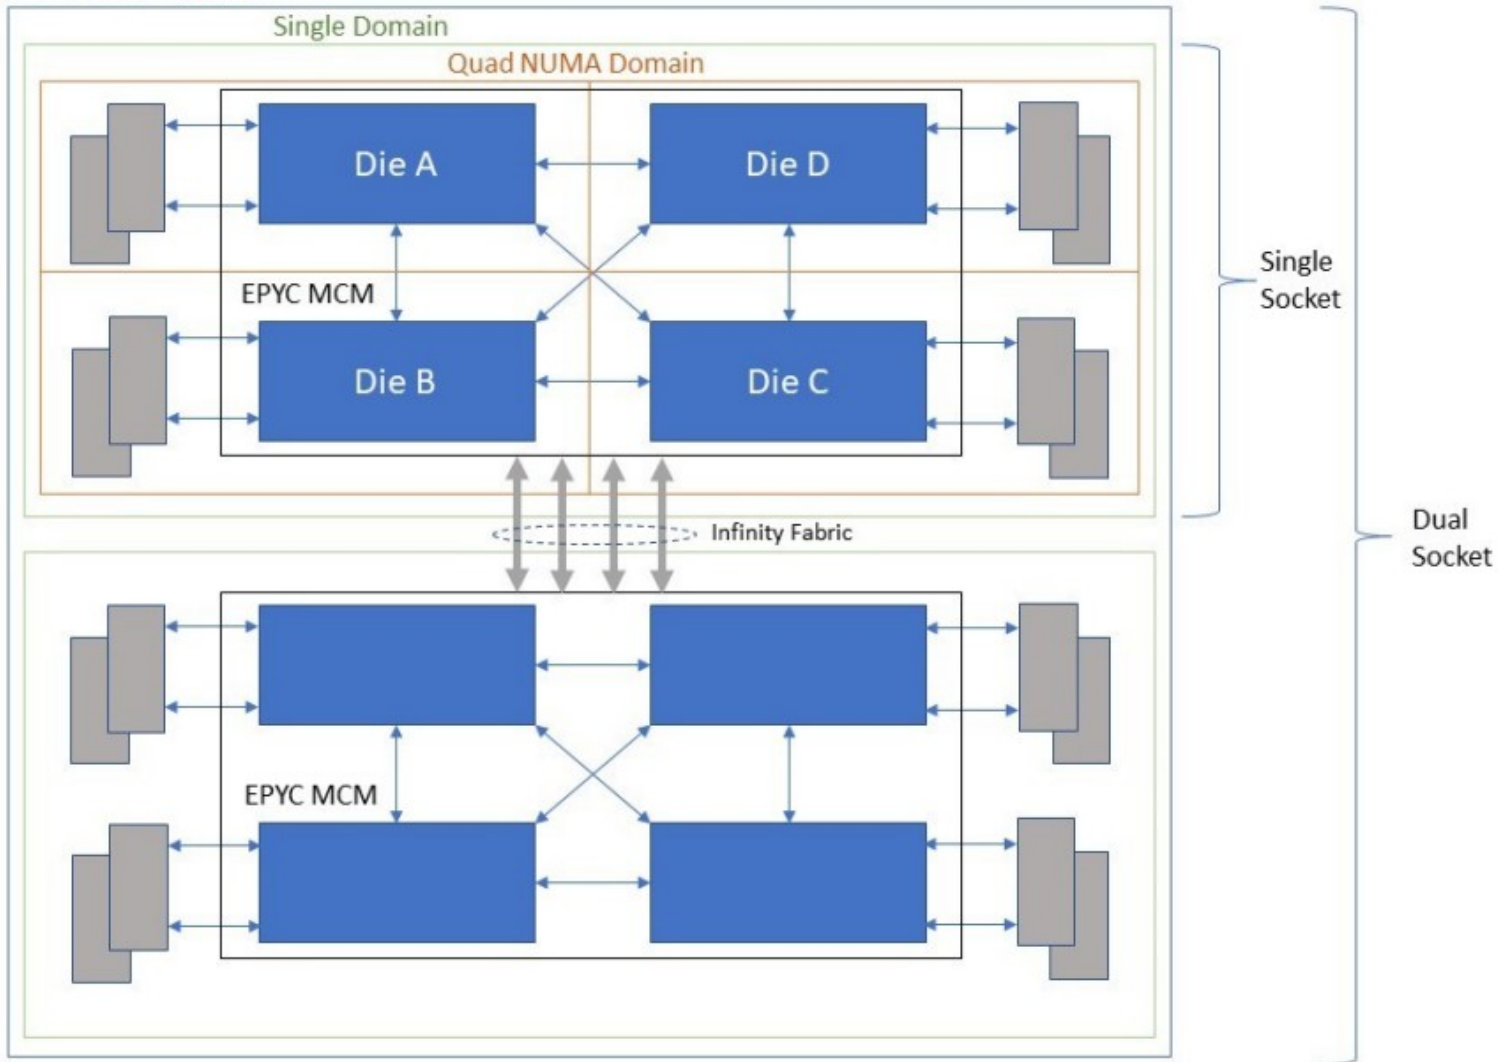
1.0X

EPYC MCM



4 x 8C Die Cost

0.59X¹



NUMA: AMD Epyc

OpenMP

Compiler notes: Other

- Linux and OS X with gcc:

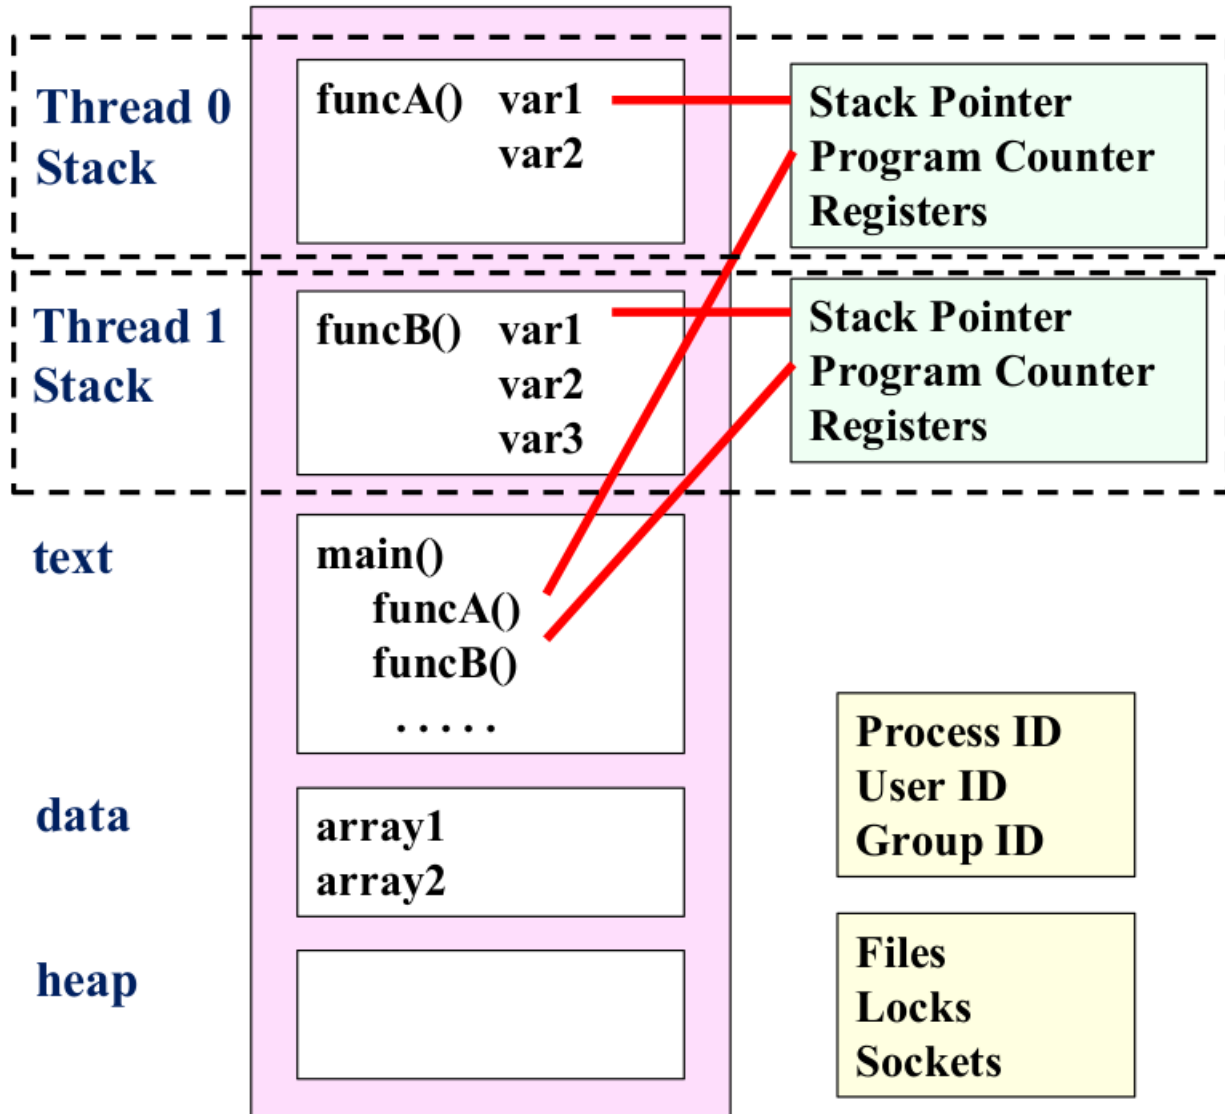
- > `gcc -fopenmp foo.c`

- > `export OMP_NUM_THREADS=4`

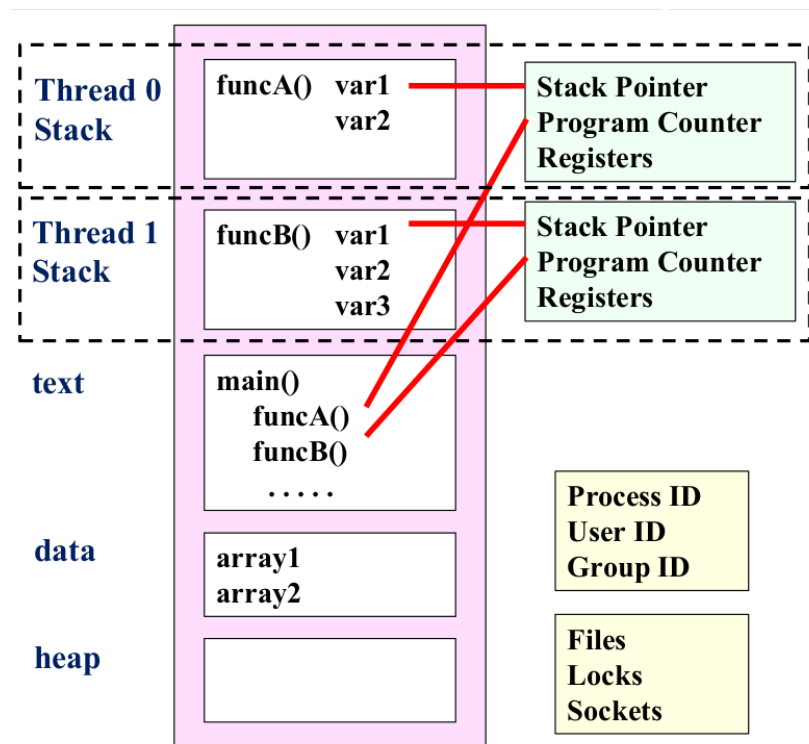


- > `./a.out`

for the Bash shell

OpenMP & NUMA






OpenMP & NUMA





Example: Medium-grained loop parallelism

Number of systems versus core count



P1	<pre>do i=1,500 a(i)=c*b(i) enddo</pre>	<pre>do i=1,1000 a(i)=c*b(i) enddo</pre>
P2	<pre>do i=501,1000 a(i)=c*b(i) enddo</pre>	

Figure 5.1: An example for medium-grained parallelism: The iterations of a loop are distributed to two processors P1 and P2 (in shared memory) for concurrent execution.

5.2.2 Functional parallelism

Example: Master-worker scheme

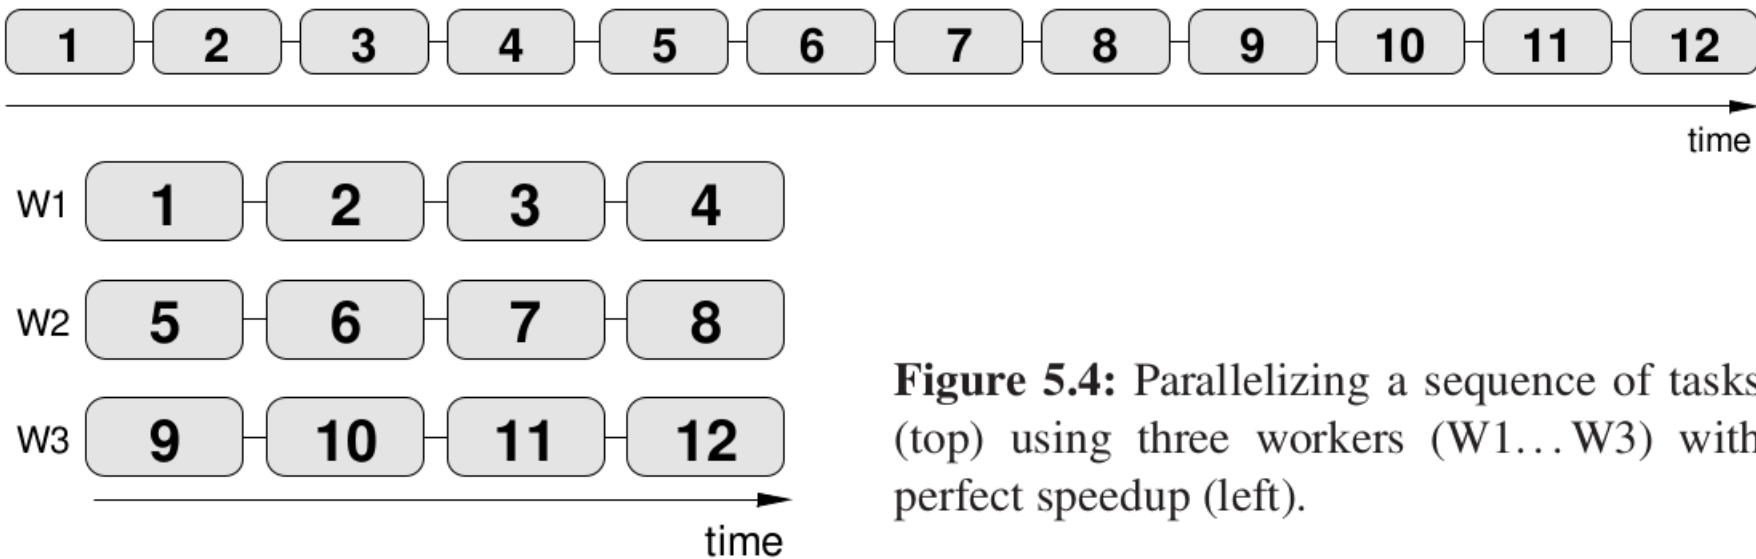


Figure 5.4: Parallelizing a sequence of tasks (top) using three workers (W1...W3) with perfect speedup (left).

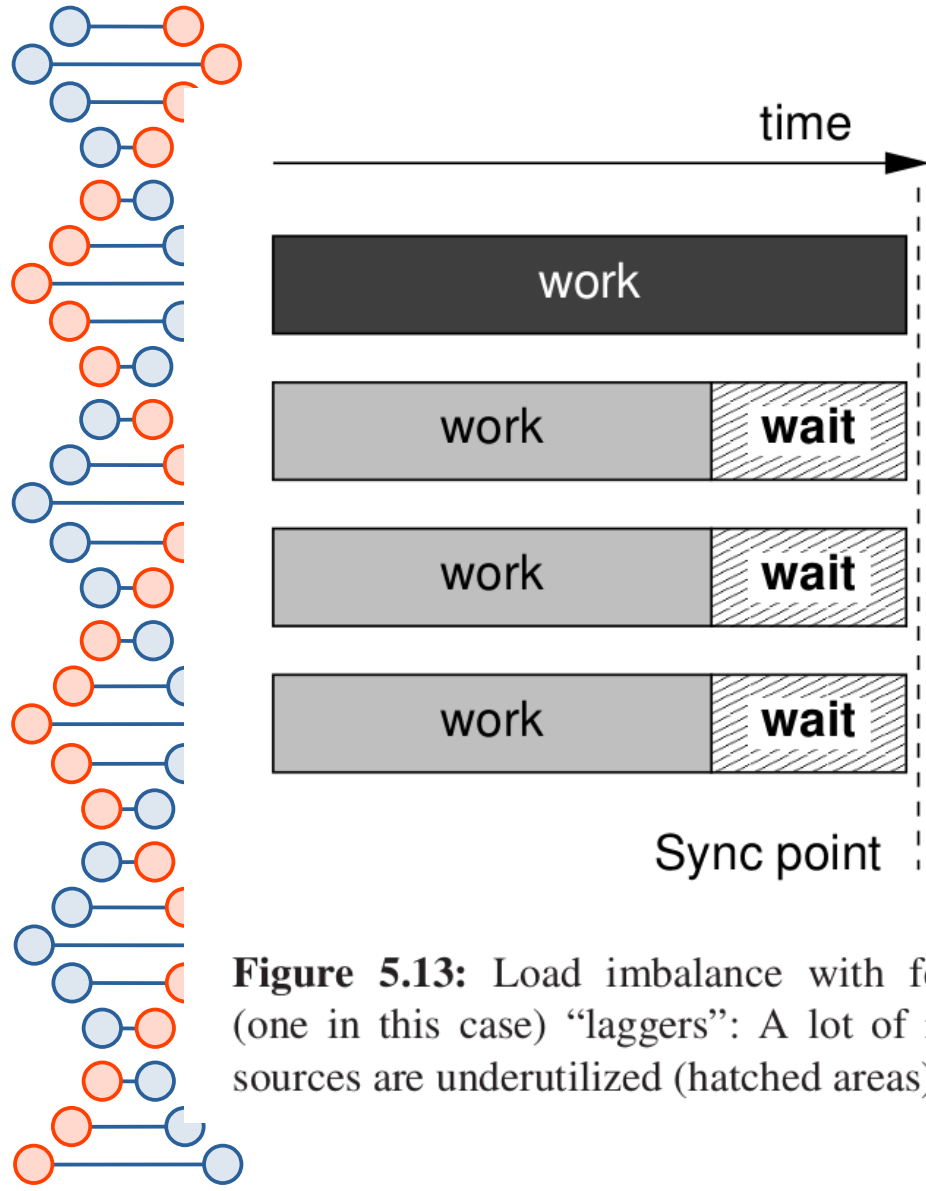


Figure 5.13: Load imbalance with few (one in this case) "laggers": A lot of resources are underutilized (hatched areas).


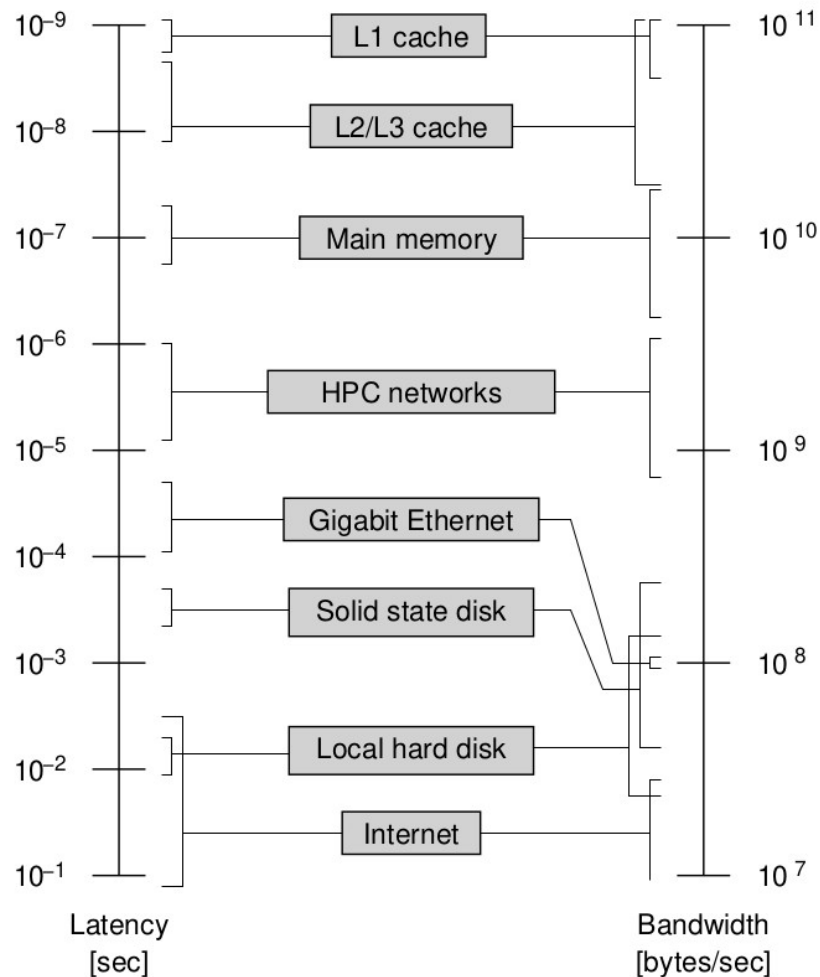



Figure 5.14: Load imbalance with few (one in this case) "speeders": Underutilization may be acceptable.

Latency vs Bandwidth

$$B_m = \frac{\text{memory bandwidth [GWords/sec]}}{\text{peak performance [GFlops/sec]}} = \frac{b_{\max}}{P_{\max}}$$






Latency vs Bandwidth

data path	balance [W/F]
cache	0.5–1.0
machine (memory)	0.03–0.5
interconnect (high speed)	0.001–0.02
interconnect (GBit ethernet)	0.0001–0.0007
disk (or disk subsystem)	0.0001–0.01

Table 3.1: Typical balance values for operations limited by different transfer paths. In case of network and disk connections, the peak performance of typical dual-socket compute nodes was taken as a basis.



1 CPU core

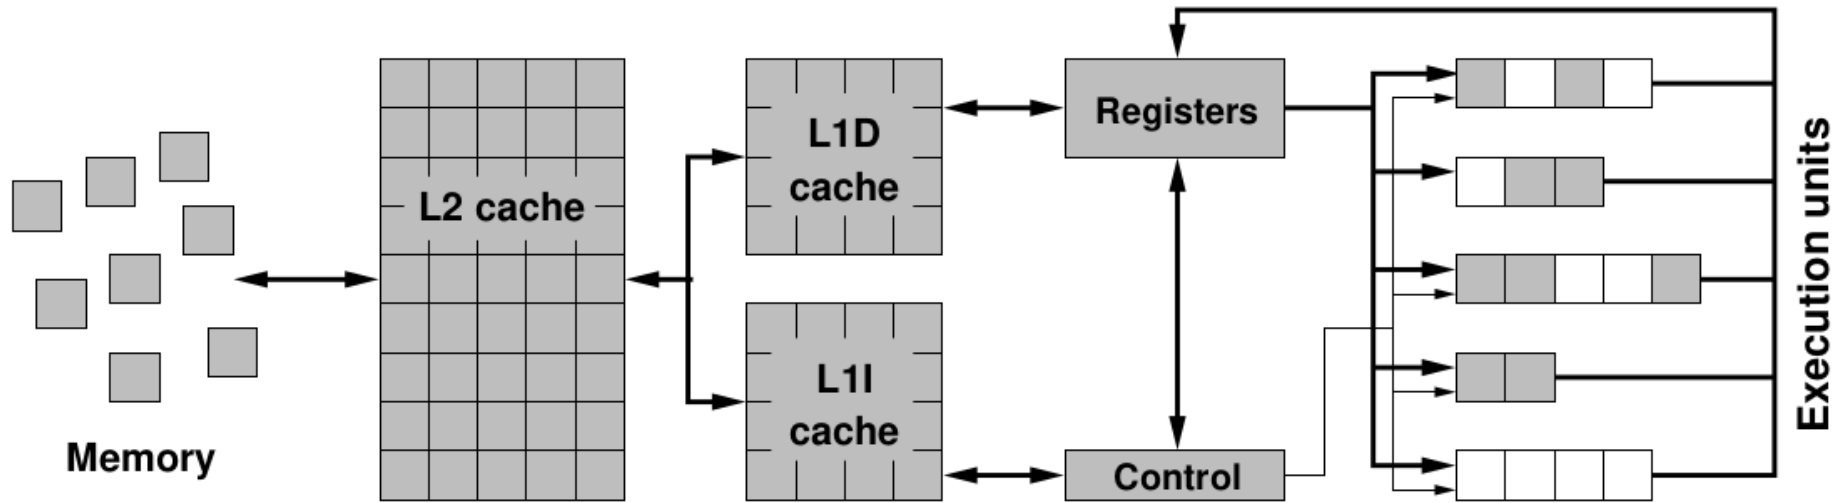


Figure 1.19: Simplified diagram of control/data flow in a (multi-)pipelined microprocessor without SMT. White boxes in the execution units denote pipeline bubbles (stall cycles). Graphics by courtesy of Intel.



Hardware performance counters Per Core

1	CPU Cycles.....	8721026107
2	Retired Instructions.....	21036052778
3	Average number of retired instructions per cycle.....	2.398151
4	L2 Misses.....	101822
5	Bus Memory Transactions.....	54413
6	Average MB/s requested by L2.....	2.241689
7	Average Bus Bandwidth (MB/s).....	1.197943
8	Retired Loads.....	694058538
9	Retired Stores.....	199529719
10	Retired FP Operations.....	7134186664
11	Average MFLOP/s.....	1225.702566
12	Full Pipe Bubbles in Main Pipe.....	3565110974
13	Percent stall/bubble cycles.....	40.642963

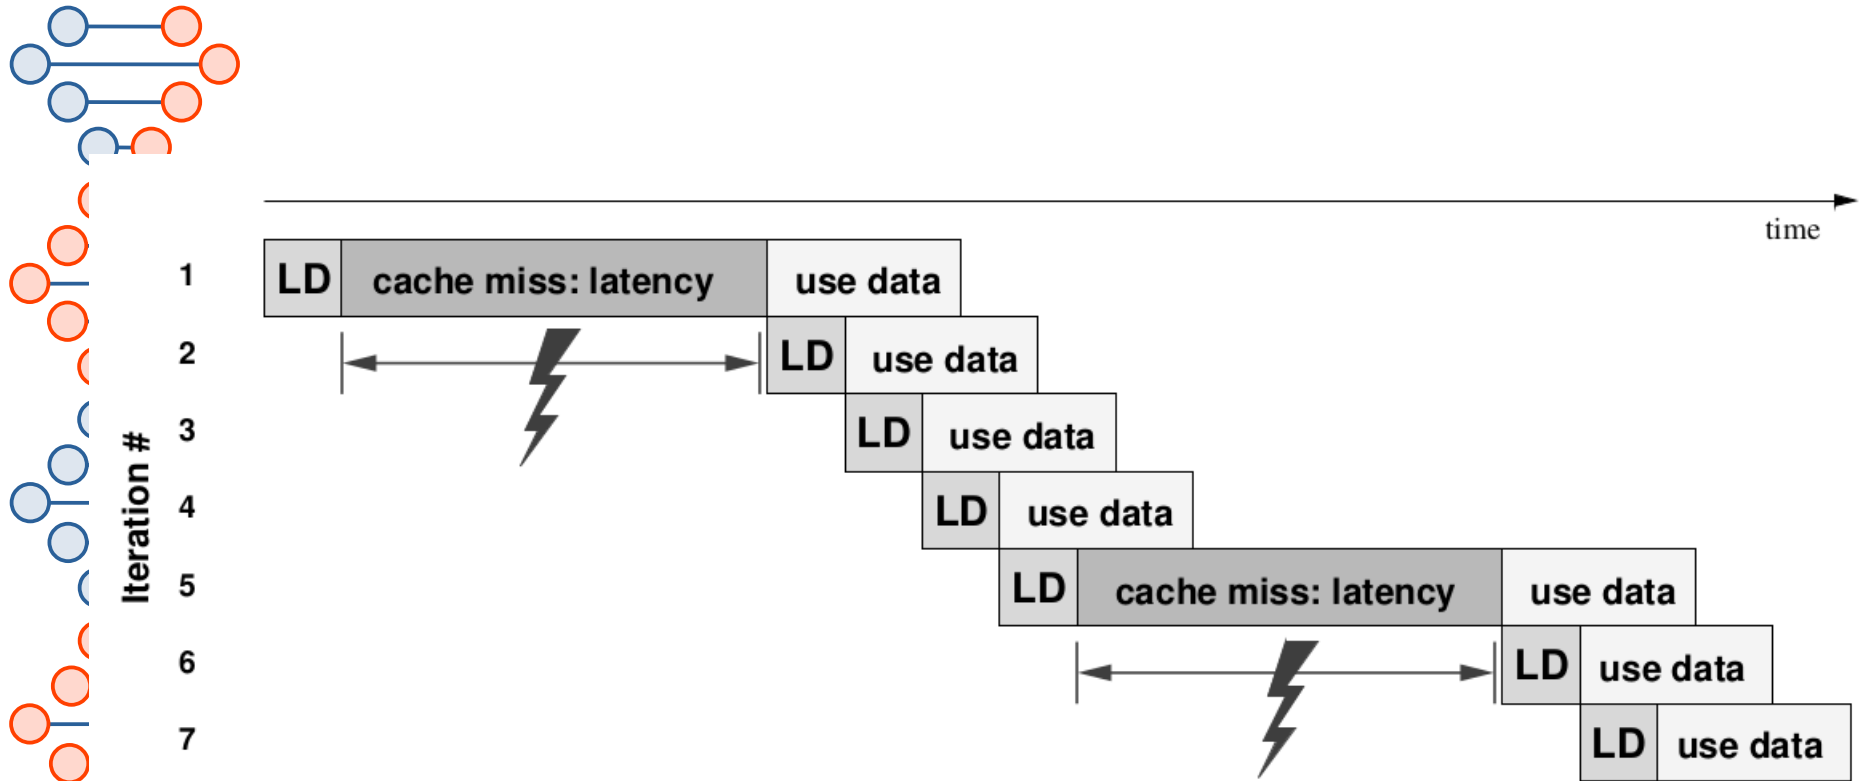


Figure 1.12: Timing diagram on the influence of cache misses and subsequent latency penalties for a vector norm loop. The penalty occurs on each new miss.

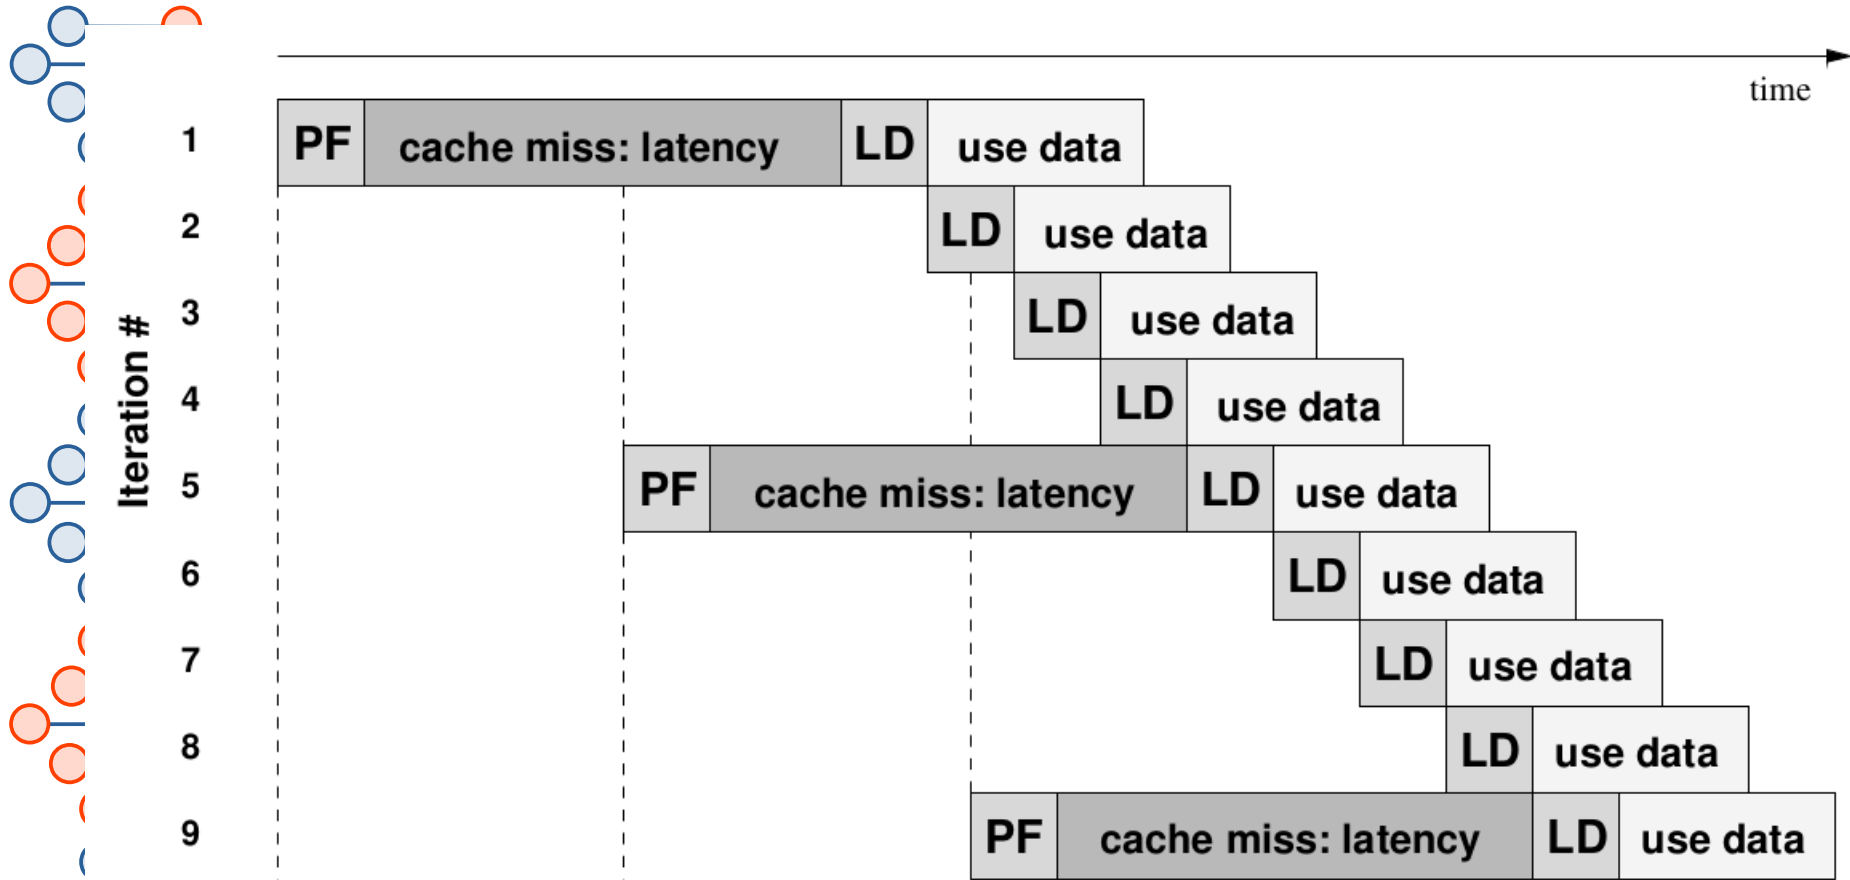



Figure 1.13: Computation and data transfer can be overlapped much better with prefetching. In this example, two outstanding prefetches are required to hide latency completely.



4.5 Networks


Point-to-point connections

Whatever the underlying hardware may be, the communication characteristics of a single point-to-point connection can usually be described by a simple model: Assuming that the total transfer time for a message of size N [bytes] is composed of latency and streaming parts,

$$T = T_\ell + \frac{N}{B} \quad (4.1)$$

and B being the maximum (asymptotic) network bandwidth in MBytes/sec, the effective bandwidth is

$$B_{\text{eff}} = \frac{N}{T_\ell + \frac{N}{B}} \quad (4.2)$$



4.5 Networks




Figure 4.10: Fit of the model for effective bandwidth (4.2) to data measured on a GigE network. The fit cannot accurately reproduce the measured value of T_ℓ (see text). $N_{1/2}$ is the message length at which half of the saturation bandwidth is reached (dashed line).

4.5 Networks




Figure 4.11: Fits of the model for effective bandwidth (4.2) to data measured on a DDR InfiniBand network. “Good” fits for asymptotic bandwidth (dotted-dashed) and latency (dashed) are shown separately, together with a fit function that unifies both (solid).

4.5 Networks





Figure 4.10: Fit of the model for effective bandwidth (4.2) to data measured on a GigE network. The fit cannot accurately reproduce the measured value of T_ℓ (see text). $N_{1/2}$ is message length at which half of the saturation bandwidth is reached (dashed line).




Figure 4.11: Fits of the model for effective bandwidth (4.2) to data measured on a DDR InfiniBand network. “Good” fits for asymptotic bandwidth (dotted-dashed) and latency (dashed) are shown separately, together with a fit function that unifies both (solid).